

Source:

1 | Part A

- What are the best and worst possible scores it could get?
 - When using f1 score, the best possible score would be 1, and the worst would be 0.
- Why do you think the model is getting the score you're seeing?
- How do the scores between the training and testing sets compare?

2 | Part B

Please answer the following questions and briefly explain your answer:

- Thinking about the R^2 metric used for evaluating regression, answer the following questions:
 - What is the best highest possible score you could get?
 - The highest possible score is 1, representing a perfect fit. In order to achieve an R^2 of 1, the error must be 0.
 - If your model simply predicted the average value of the training set no matter what the input was, what score would you get on a test set whose average matched that of the training set?
 - The score would be 0, as the denominator and numerator would be equal. This would result in the equation $1-1$, which equals 0.
 - What is the lowest score that you can get?
 - Negative infinity.
- When using accuracy to measure your model's performance on a classification problem:
 - What is the best possible score you could get?
 - 1. In a scenario with x samples, an entirely correct classification would lead to x/x , which is 1.
 - If your model always predicted the same class no matter what the input, what score would you get on a test set where 85% of the items were in that class?
 - 85%
- What is the worst possible score you can get on a dataset that only has two classes?
 - 0. Every sample could be sorted incorrectly as long as there is more than one class.
- A model gets a recall score of 0 for class A on a test set with classes A, B, and C. If you take one of the test items that is in class A and have this model predict what class it is, what will it predict?
 - A recall score of 0 for class A means that every item in class A was labeled incorrectly. When given another sample in class A, the model will predict either class B or class C.
- A model gets a precision score of 1 for class A on a test set with classes A, B, and C. If you take one of the test items that is in class A and have your model predict what class it is, what will it predict?
 - The model could predict any of the three classes. A precision value is only effected by true positive and false positive rate, meaning that when given a sample in class A, predicting C or B would not lower class A's precision.

- If a model with classes A and B has an AUC score of 1 and you give it an item from the test set that is in class A, what class will it predict and what probability will it give for that class?
 - It will predict class A, and give a probability of 100%. This is because an AUC score of 1 represents a perfect model.
- If a model with classes A and B has an AUC score of 0 and you give it an item from the test set that is in class A, what class will it predict and what probability will it give for that class?
 - This means that the model is always perfectly incorrect. If you were to give it a sample from class B, it will always pick class A, and vice versa. The model would still give a probability of 100%.